

# **CHAPTER 26**

## **Enhanced Data Models: Introduction to Active, Temporal, Spatial, Multimedia, and Deductive Databases**

# 26.1 Active Database Concepts and Triggers

- Database systems implement rules that specify actions automatically triggered by certain events
- Triggers
  - Technique for specifying certain types of active rules
- Commercial relational DBMSs have various versions of triggers available
  - Oracle syntax used to illustrate concepts

# Generalized Model for Active Databases and Oracle Triggers

- Event-condition-action (ECA) model
  - Event triggers a rule
    - Usually database update operations
  - Condition determines whether rule action should be completed
    - Optional
    - Action will complete only if condition evaluates to true
  - Action to be taken
    - Sequence of SQL statements, transaction, or external program

# Example

- Events that may cause a change in value of Total\_sal attribute
  - Inserting new employee
  - Changing salary
  - Reassigning or deleting employees

## EMPLOYEE

Name	<u>Ssn</u>	Salary	Dno	Supervisor_ssn
------	------------	--------	-----	----------------

## DEPARTMENT

Dname	<u>Dno</u>	Total_sal	Manager_ssn
-------	------------	-----------	-------------

Figure 26.1 A simplified COMPANY database used for active rule examples

# Example (cont'd.)

- Condition to be evaluated
  - Check that value of Dno attribute is not NULL
- Action to be taken
  - Automatically update the value of Total\_sal

```

(a) R1: CREATE TRIGGER Total_sal1
        AFTER INSERT ON EMPLOYEE
        FOR EACH ROW
        WHEN ( NEW.Dno IS NOT NULL )
            UPDATE DEPARTMENT
            SET Total_sal = Total_sal + NEW.Salary
            WHERE Dno = NEW.Dno;

R2: CREATE TRIGGER Total_sal2
        AFTER UPDATE OF Salary ON EMPLOYEE
        FOR EACH ROW
        WHEN ( NEW.Dno IS NOT NULL )
            UPDATE DEPARTMENT
            SET Total_sal = Total_sal + NEW.Salary - OLD.Salary
            WHERE Dno = NEW.Dno;

R3: CREATE TRIGGER Total_sal3
        AFTER UPDATE OF Dno ON EMPLOYEE
        FOR EACH ROW
        BEGIN
            UPDATE DEPARTMENT
            SET Total_sal = Total_sal + NEW.Salary
            WHERE Dno = NEW.Dno;
            UPDATE DEPARTMENT
            SET Total_sal = Total_sal - OLD.Salary
            WHERE Dno = OLD.Dno;
        END;

```

Figure 26.2 Specifying active rules as triggers in Oracle notation (a) Triggers for automatically maintaining the consistency of Total\_sal of DEPARTMENT

```
R4: CREATE TRIGGER Total_sal4
AFTER DELETE ON EMPLOYEE
FOR EACH ROW
WHEN ( OLD.Dno IS NOT NULL)
UPDATE DEPARTMENT
SET Total_sal = Total_sal - OLD.Salary
WHERE Dno = OLD.Dno;
```

```
(b) R5: CREATE TRIGGER Inform_supervisor1
BEFORE INSERT OR UPDATE OF Salary, Supervisor_ssn
ON EMPLOYEE
FOR EACH ROW
WHEN ( NEW.Salary > ( SELECT Salary FROM EMPLOYEE
                      WHERE Ssn = NEW.Supervisor_ssn ) )
inform_supervisor(NEW.Supervisor_ssn, NEW.Ssn );
```

Figure 26.2 (cont'd.) Specifying active rules as triggers in Oracle notation (b)  
Trigger for comparing an employee's salary with that of his or her supervisor



# Design and Implementation Issues for Active Databases

- Deactivated rule
  - Will not be triggered by the triggering event
- Activate command
  - Makes the rule active again
- Drop command
  - Deletes the rule from the system
- Approach: group rules into rule sets
  - Entire rule set can be activated, deactivated, or dropped

# Design and Implementation Issues for Active Databases (cont'd.)

- Timing of action
  - Before trigger executes trigger before executing event that caused the trigger
  - After trigger executes trigger after executing the event
  - Instead of trigger executes trigger instead of executing the event
- Action can be considered separate transaction
  - Or part of same transaction that triggered the rule

# Design and Implementation Issues for Active Databases (cont'd.)

- Rule consideration
  - Immediate consideration
    - Condition evaluated as part of same transaction
    - Evaluate condition either before, after, or instead of executing the triggering event
  - Deferred consideration
    - Condition evaluated at the end of the transaction
  - Detached consideration
    - Condition evaluated as a separate transaction

# Design and Implementation Issues for Active Databases (cont'd.)

- Row-level rule
  - Rule considered separately for each row
- Statement-level rule
  - Rule considered once for entire statement
- Difficult to guarantee consistency and termination of rules

# Examples of Statement-Level Active Rules in STARBURST

```
R1S: CREATE RULE Total_sal1 ON EMPLOYEE
      WHEN   INSERTED
      IF     EXISTS      ( SELECT * FROM INSERTED WHERE Dno IS NOT NULL)
      THEN   UPDATE      DEPARTMENT AS D
            SET          D.Total_sal = D.Total_sal +
                        ( SELECT SUM (I.Salary) FROM INSERTED AS I WHERE D.Dno = I.Dno )
            WHERE        D.Dno IN ( SELECT Dno FROM INSERTED );

R2S: CREATE RULE Total_sal2 ON EMPLOYEE
      WHEN   UPDATED    ( Salary )
      IF     EXISTS      ( SELECT * FROM NEW-UPDATED WHERE Dno IS NOT NULL)
            OR EXISTS    ( SELECT * FROM OLD-UPDATED WHERE Dno IS NOT NULL)
      THEN   UPDATE      DEPARTMENT AS D
            SET          D.Total_sal = D.Total_sal +
                        ( SELECT SUM (N.Salary) FROM NEW-UPDATED AS N
                          WHERE D.Dno = N.Dno ) -
                        ( SELECT SUM (O.Salary) FROM OLD-UPDATED AS O
                          WHERE D.Dno = O.Dno )
            WHERE        D.Dno IN ( SELECT Dno FROM NEW-UPDATED ) OR
                        D.Dno IN ( SELECT Dno FROM OLD-UPDATED );
```

Figure 26.5 (continues) Active rules using statement-level semantics in STARBURST notation

# Examples of Statement-Level Active Rules in STARBURST (cont'd.)

```
R3S: CREATE RULE Total_sal3 ON EMPLOYEE
      WHEN   UPDATED   ( Dno )
      THEN   UPDATE    DEPARTMENT AS D
              SET       D.Total_sal = D.Total_sal +
                        ( SELECT SUM (N.Salary) FROM NEW-UPDATED AS N
                          WHERE D.Dno = N.Dno )
      WHERE   D.Dno IN ( SELECT Dno FROM NEW-UPDATED );
      UPDATE  DEPARTMENT AS D
      SET     D.Total_sal = Total_sal -
            ( SELECT SUM (O.Salary) FROM OLD-UPDATED AS O
              WHERE D.Dno = O.Dno )
      WHERE   D.Dno IN ( SELECT Dno FROM OLD-UPDATED );
```

Figure 26.5 (cont'd.) Active rules using statement-level semantics in STARBURST notation

# Potential Applications for Active Databases

- Allow notification of certain conditions that occur
- Enforce integrity constraints
- Automatically maintain derived data
- Maintain consistency of materialized views
- Enable consistency of replicated tables

# Triggers in SQL-99

```
T1: CREATE TRIGGER Total_sal1
    AFTER UPDATE OF Salary ON EMPLOYEE
    REFERENCING OLD ROW AS O, NEW ROW AS N
    FOR EACH ROW
    WHEN ( N.Dno IS NOT NULL )
    UPDATE DEPARTMENT
    SET Total_sal = Total_sal + N.salary - O.salary
    WHERE Dno = N.Dno;

T2: CREATE TRIGGER Total_sal2
    AFTER UPDATE OF Salary ON EMPLOYEE
    REFERENCING OLD TABLE AS O, NEW TABLE AS N
    FOR EACH STATEMENT
    WHEN EXISTS ( SELECT * FROM N WHERE N.Dno IS NOT NULL ) OR
        EXISTS ( SELECT * FROM O WHERE O.Dno IS NOT NULL )
    UPDATE DEPARTMENT AS D
    SET D.Total_sal = D.Total_sal
    + ( SELECT SUM (N.Salary) FROM N WHERE D.Dno=N.Dno )
    - ( SELECT SUM (O.Salary) FROM O WHERE D.Dno=O.Dno )
    WHERE Dno IN ( ( SELECT Dno FROM N ) UNION ( SELECT Dno FROM O ) );
```

Figure 26.6 Trigger T1 illustrating the syntax for defining triggers in SQL-99



## 26.2 Temporal Database Concepts

- Temporal databases require some aspect of time when organizing information
  - Healthcare
  - Insurance
  - Reservation systems
  - Scientific databases
- Time considered as ordered sequence of points
  - Granularity determined by the application

# Temporal Database Concepts (cont'd.)

- Chronon
  - Term used to describe minimal granularity of a particular application
- Reference point for measuring specific time events
  - Various calendars
- SQL2 temporal data types
  - DATE, TIME, TIMESTAMP, INTERVAL, PERIOD

# Temporal Database Concepts (cont'd.)

- Point events or facts
  - Typically associated with a single time point
  - Time series data
- Duration events or facts
  - Associated with specific time period
  - Time period represented by start and end points
- Valid time
  - True in the real world

# Temporal Database Concepts (cont'd.)

- Transaction time
  - Value of the system clock when information is valid in the system
- User-defined time
- Bitemporal database
  - Uses valid time and transaction time
- Valid time relations
  - Used to represent history of changes

# Temporal Database Concepts (cont'd.)

(a) EMP\_VT

Name	<u>Ssn</u>	Salary	Dno	Supervisor_ssn	<u>Vst</u>	Vet
------	------------	--------	-----	----------------	------------	-----

DEPT\_VT

Dname	<u>Dno</u>	Total_sal	Manager_ssn	<u>Vst</u>	Vet
-------	------------	-----------	-------------	------------	-----

(b) EMP\_TT

Name	<u>Ssn</u>	Salary	Dno	Supervisor_ssn	<u>Tst</u>	Tet
------	------------	--------	-----	----------------	------------	-----

DEPT\_TT

Dname	<u>Dno</u>	Total_sal	Manager_ssn	<u>Tst</u>	Tet
-------	------------	-----------	-------------	------------	-----

(c) EMP\_BT

Name	<u>Ssn</u>	Salary	Dno	Supervisor_ssn	<u>Vst</u>	Vet	<u>Tst</u>	Tet
------	------------	--------	-----	----------------	------------	-----	------------	-----

DEPT\_BT

Dname	<u>Dno</u>	Total_sal	Manager_ssn	<u>Vst</u>	Vet	<u>Tst</u>	Tet
-------	------------	-----------	-------------	------------	-----	------------	-----

Figure 26.7 Different types of temporal relational databases (a) Valid time database schema (b) Transaction time database schema (c) Bitemporal database schema

# Temporal Database Concepts (cont'd.)

EMP\_VT

Name	<u>Ssn</u>	Salary	Dno	Supervisor_ssn	<u>Vst</u>	Vet
Smith	123456789	25000	5	333445555	2002-06-15	2003-05-31
Smith	123456789	30000	5	333445555	2003-06-01	Now
Wong	333445555	25000	4	999887777	1999-08-20	2001-01-31
Wong	333445555	30000	5	999887777	2001-02-01	2002-03-31
Wong	333445555	40000	5	888665555	2002-04-01	Now
Brown	222447777	28000	4	999887777	2001-05-01	2002-08-10
Narayan	666884444	38000	5	333445555	2003-08-01	Now

...

DEPT\_VT

Dname	<u>Dno</u>	Manager_ssn	<u>Vst</u>	Vet
Research	5	888665555	2001-09-20	2002-03-31
Research	5	333445555	2002-04-01	Now

...

Figure 26.8 Some tuple versions in the valid time relations EMP\_VT and DEPT\_VT

# Temporal Database Concepts (cont'd.)

- Types of updates
  - Proactive
  - Retroactive
  - Simultaneous
- Timestamp recorded whenever change is applied to database
- Bitemporal relations
  - Application requires both valid time and transaction time

# Temporal Database Concepts (cont'd.)

- Implementation considerations
  - Store all tuples in the same table
  - Create two tables: one for currently valid information and one for the rest
  - Vertically partition temporal relation attributes into separate relations
    - New tuple created whenever any attribute updated
- Append-only database
  - Keeps complete record of changes and corrections



# Temporal Database Concepts (cont'd.)

- Attribute versioning
  - Simple complex object used to store all temporal changes of the object
  - Time-varying attribute
    - Values versioned over time by adding temporal periods to the attribute
  - Non-time-varying attribute
    - Values do not change over time

```

class TEMPORAL_SALARY
{   attribute   Date           Valid_start_time;
    attribute   Date           Valid_end_time;
    attribute   float          Salary;
};

class TEMPORAL_DEPT
{   attribute   Date           Valid_start_time;
    attribute   Date           Valid_end_time;
    attribute   DEPARTMENT_VT  Dept;
};

class TEMPORAL_SUPERVISOR
{   attribute   Date           Valid_start_time;
    attribute   Date           Valid_end_time;
    attribute   EMPLOYEE_VT    Supervisor;
};

class TEMPORAL_LIFESPAN
{   attribute   Date           Valid_start time;
    attribute   Date           Valid end time;
};

class EMPLOYEE_VT
(   extent EMPLOYEES )
{   attribute   list<TEMPORAL_LIFESPAN>    lifespan;
    attribute   string                     Name;
    attribute   string                     Ssn;
    attribute   list<TEMPORAL_SALARY>      Sal_history;
    attribute   list<TEMPORAL_DEPT>       Dept_history;
    attribute   list <TEMPORAL_SUPERVISOR> Supervisor_history;
};

```

Figure 26.10 Possible ODL schema for a temporal valid time EMPLOYEE\_VT object class using attribute versioning

# Temporal Database Concepts (cont'd.)

- TSQL2 language
  - Extends SQL for querying valid time and transaction time tables
  - Used to specify whether a relation is temporal or nontemporal
- Temporal database query conditions may involve time and attributes
  - Pure time condition involves only time
  - Attribute and time conditions

# Temporal Database Concepts (cont'd.)

- CREATE TABLE statement
  - Extended with optional AS clause
  - Allows users to declare different temporal options
  - Examples:
    - AS VALID STATE<GRANULARITY> (valid time relation with valid time period)
    - AS TRANSACTION (transaction time relation with transaction time period)
- Keywords STATE and EVENT
  - Specify whether a time period or point is associated with valid time dimension

# Temporal Database Concepts (cont'd.)

- Time series data
  - Often used in financial, sales, and economics applications
  - Special type of valid event data
  - Event's time points predetermined according to fixed calendar
  - Managed using specialized time series management systems
  - Supported by some commercial DBMS packages

## 26.3 Spatial Database Concepts

- Spatial databases support information about objects in multidimensional space
  - Examples: cartographic databases, geographic information systems, weather information databases
- Spatial relationships among the objects are important
- Optimized to query data such as points, lines, and polygons
  - Spatial queries

# Spatial Database Concepts (cont'd.)

- Measurement operations
  - Used to measure global properties of single objects
- Spatial analysis operations
  - Uncover spatial relationships within and among mapped data layers
- Flow analysis operations
  - Help determine shortest path between two points

# Spatial Database Concepts (cont'd.)

- Location analysis
  - Determine whether given set of points and lines lie within a given polygon
- Digital terrain analysis
  - Used to build three-dimensional models



# Spatial Database Concepts (cont'd.)

Analysis Type	Type of Operations and Measurements
Measurements	Distance, perimeter, shape, adjacency, and direction
Spatial analysis/statistics	Pattern, autocorrelation, and indexes of similarity and topology using spatial and nonspatial data
Flow analysis	Connectivity and shortest path
Location analysis	Analysis of points and lines within a polygon
Terrain analysis	Slope/aspect, catchment area, drainage network
Search	Thematic search, search by region

Table 26.1 Common types of analysis for spatial data

# Spatial Database Concepts (cont'd.)

- Spatial data types
  - Map data
    - Geographic or spatial features of objects in a map
  - Attribute data
    - Descriptive data associated with map features
  - Image data
    - Satellite images
- Models of spatial information
  - Field models
  - Object models

# Spatial Database Concepts (cont'd.)

- Spatial operator categories
  - Topological operators
    - Properties do not change when topological transformations applied
  - Projective operators
    - Express concavity/convexity of objects
  - Metric operators
    - Specifically describe object's geometry
  - Dynamic spatial operators
    - Create, destroy, and update

# Spatial Database Concepts (cont'd.)

- Spatial queries
  - Range queries
    - Example: find all hospitals with the Metropolitan Atlanta city area
  - Nearest neighbor queries
    - Example: find police car nearest location of a crime
  - Spatial joins or overlays
    - Example: find all homes within two miles of a lake

# Spatial Database Concepts (cont'd.)

- Spatial data indexing
  - Grid files
  - R-trees
  - Spatial join index
- Spatial data mining techniques
  - Spatial classification
  - Spatial association
  - Spatial clustering

# 26.4 Multimedia Database Concepts

- Multimedia databases allow users to store and query images, video, audio, and documents
- Content-based retrieval
  - Automatic analysis
  - Manual identification
  - Color often used in content-based image retrieval
  - Texture and shape
- Object recognition
  - Scale-invariant feature transform (SIFT) approach

# Multimedia Database Concepts (cont'd.)

- Semantic tagging of images
  - User-supplied tags
  - Automated generation of image tags
  - Web Ontology Language (OWL) provides concept hierarchy
- Analysis of audio data sources
  - Text-based indexing
  - Content-based indexing

# 26.5 Introduction to Deductive Databases

- Deductive database uses facts and rules
  - Inference engine can deduce new facts using rules
- Prolog/Datalog notation
  - Based on providing predicates with unique names
  - Predicate has an implicit meaning and a fixed number of arguments
    - If arguments are all constant values, predicate states that a certain fact is true
    - If arguments are variables, considered as a query or part of a rule or constraint



# Prolog Notation and The Supervisory Tree

(a)

## Facts

```
SUPERVISE(franklin, john).  
SUPERVISE(franklin, ramesh).  
SUPERVISE(franklin, joyce).  
SUPERVISE(jennifer, alicia).  
SUPERVISE(jennifer, ahmad).  
SUPERVISE(james, franklin).  
SUPERVISE(james, jennifer).  
...
```

## Rules

```
SUPERIOR(X, Y) :- SUPERVISE(X, Y).  
SUPERIOR(X, Y) :- SUPERVISE(X, Z), SUPERIOR(Z, Y).  
SUBORDINATE(X, Y) :- SUPERIOR(Y, X).
```

## Queries

```
SUPERIOR(james, Y)?  
SUPERIOR(james, joyce)?
```

(b)




Figure 26.11 (a) Prolog notation (b) The supervisory tree

# Introduction to Deductive Databases (cont'd.)

- Datalog notation
  - Program built from basic objects called atomic formulas
  - Literals of the form  $p(a_1, a_2, \dots, a_n)$ 
    - $p$  is the predicate name
    - $n$  is the number of arguments for predicate  $p$
- Interpretations of rules
  - Proof-theoretic versus model-theoretic
  - Deductive axioms
  - Ground axioms

# Introduction to Deductive Databases (cont'd.)

- |   |                           |
|---|---------------------------|
| 1. SUPERIOR(X, Y) :- SUPERVISE(X, Y).                 | (rule 1)                  |
| 2. SUPERIOR(X, Y) :- SUPERVISE(X, Z), SUPERIOR(Z, Y). | (rule 2)                  |
| 3. SUPERVISE(jennifer, ahmad).                        | (ground axiom, given)     |
| 4. SUPERVISE(james, jennifer).                        | (ground axiom, given)     |
| 5. SUPERIOR(jennifer, ahmad).                         | (apply rule 1 on 3)       |
| 6. SUPERIOR(james, ahmad).                            | (apply rule 2 on 4 and 5) |

Figure 26.12 Proving a new fact

# Introduction to Deductive Databases (cont'd.)

- Safe program or rule
  - Generates a finite set of facts
- Nonrecursive query
  - Includes only nonrecursive predicates

# Use of Relational Operations

Figure 26.16  
Predicates for  
illustrating  
relational operations

```
REL_ONE(A, B, C).
REL_TWO(D, E, F).
REL_THREE(G, H, I, J).

SELECT_ONE_A_EQ_C(X, Y, Z) :- REL_ONE(C, Y, Z).
SELECT_ONE_B_LESS_5(X, Y, Z) :- REL_ONE(X, Y, Z), Y < 5.
SELECT_ONE_A_EQ_C_AND_B_LESS_5(X, Y, Z) :- REL_ONE(C, Y, Z), Y < 5.

SELECT_ONE_A_EQ_C_OR_B_LESS_5(X, Y, Z) :- REL_ONE(C, Y, Z).
SELECT_ONE_A_EQ_C_OR_B_LESS_5(X, Y, Z) :- REL_ONE(X, Y, Z), Y < 5.

PROJECT_THREE_ON_G_H(W, X) :- REL_THREE(W, X, Y, Z).

UNION_ONE_TWO(X, Y, Z) :- REL_ONE(X, Y, Z).
UNION_ONE_TWO(X, Y, Z) :- REL_TWO(X, Y, Z).

INTERSECT_ONE_TWO(X, Y, Z) :- REL_ONE(X, Y, Z), REL_TWO(X, Y, Z).

DIFFERENCE_TWO_ONE(X, Y, Z) :- _TWO(X, Y, Z) NOT(REL_ONE(X, Y, Z)).

CART_PROD_ONE_THREE(T, U, V, W, X, Y, Z) :-
    REL_ONE(T, U, V), REL_THREE(W, X, Y, Z).

NATURAL_JOIN_ONE_THREE_C_EQ_G(U, V, W, X, Y, Z) :-
    REL_ONE(U, V, W), REL_THREE(W, X, Y, Z).
```

# 26.6 Summary

- Active databases
  - Specify active rules
- Temporal databases
  - Involve time concepts
- Spatial databases
  - Involve spatial characteristics
- Multimedia databases
  - Store images, audio, video, documents, and more
- Deductive databases
  - Prolog and Datalog notation